

# INTRODUCTION TO ETHICS

PHIL 126, YALE UNIVERSITY, SPRING 2019

These are lecture notes for PHIL 175b, “Introduction to Ethics,” taught by Shelly Kagan at Yale University during the spring of 2019. These notes are not official, and have not been proofread by the instructor for the course. They live in my lecture notes respository at

<https://github.com/jopetty/lecture-notes/tree/master/PHIL-175>.

If you find any errors, please open a bug report describing the error and label it with the course identifier, or open a pull request so I can correct it.

## Contents

1	January 15, 2019	1
2	January 17, 2019	2
3	Tuesday, January 22	4
	3.1 <i>Why have morality?</i> . . . . .	4
4	Thursday, January 24	7
	4.1 <i>What is the correct theory of the good?</i> . . . . .	7
5	Tuesday, January 29	9
6	Thursday, January 31	11
7	No Notes	12
8	Thursday, February 7	13

## 1 January 15, 2019

“You are welcome to sit adoringly at  
my feet if you so wish.”

---

Shelly

Shelly (he asked us to call him that) talked about how there really isn't a good definition of what Philosophy is, perhaps unlike many other academic disciplines. Rather, it is defined by the types of things we ask questions about and the types of thoughts we study. Perhaps it's a “Roller-coaster for your mind.”

**Definition (Ethics).** Moral Philosophy or Ethics is the attempt to answer the question “How should I live?” In the course of answering this, various other questions will arise and must be answered along the way.

*Ethics*

This doesn't deal with the empirics of “How *do* we live.” That is a realm for sociologists. We're also not really interested in people's opinions on justified actions — we're concerned with what the truth of the matter actually is. There will be no “comparative ethics” here. We're also not doing a history of philosophical ideas here.

Philosophy thrives on rival views, not on proven results. There isn't a philosophical answer to “what is well being?” — you have to decide for yourself how to answer that question.

## 2 January 17, 2019

“The blood gurgled out of his neck as he gasped his last breath.”

---

Shelly

### *Outline of Lecture*

1. Application of Morality — Applied/Practical Ethics
2. Content of Morality — Normative Ethics
3. Foundation of Morality — Normative & Metaethics
4. Nature of Morality — Normative and Metaethics

### *Questions of Moral Philosophy*

**Example 2.1** (Questions of Moral Philosophy).

- (a) What are the various moral rules? What are the specific details? For example, if “Don’t kill” is the rule, is it every permissible to kill people? When and why? If a child picks up a loaded gun and the only way to stop them from killing people is to shoot the child, is that moral? (*This is called an innocent threat*) Are there proportionality considerations?
- (b) Do we have obligations to keep our promises? What if we have conflicting promises? How should we weight our promises? Are there ever a valid set of rules we can follow? Can we break our promises if a more pressing situation arises where we could do more good (saving someone in cardiac arrest but skipping a meeting to do so)?
- (c) Can war ever be moral? What about collateral damage and civilian deaths?
- (d) Is abortion moral or immoral? Does killing a fetus break the “don’t kill” command? When do we have rights to life?

The subfields of morality (although there really isn’t a clear line between these areas):

1. **Content of Morality** — What are the rules of a just society? What factors make an action right or wrong?

2. **Application of Morality** — How can we apply the rules of morality to real-life situations? What happens when situations are ambiguous? What conflicts arise? Abortion and Capital Punishment? This class doesn't really touch on this.
3. **Foundation of Morality** — Can our rules be derived from one another? Is a prohibition on killing entirely separate from "Do not harm people" or is it just a specific case one the more general rule? How to we systematize ethical rules? Is there a single basic rule from which all others derive (like the Golden Rule or Categorical Imperative or Utilitarianism)? What makes the true rules true? Is true even the right word? Why are moral systems just? By virtue of what is morality founded? Divine commandment, contractarianism, etc.
4. **Nature of Morality** — What's ethics all about? What's the point of ethics? What does it even mean to be moral? Why are some things more valid than others? Why do values even exist? How do we include morality in a scientific worldview? Are there "morality facts"? How could we ever hope to know whether something is right or wrong? What do "right" and "wrong" even mean? We probably won't touch on this here either (Take Moral Skepticism with Shelly in the fall).

### 3 Tuesday, January 22

“If you planted a bomb in a school yard,  
I guess that would be a pretty bad  
thing”

---

Shelly

A lot of what drives moral philosophy is “why be moral,” but to answer this we first must have an idea of what morality actually is. Our first framework will be *Utilitarianism*, popularized by Mill and Bentham, which we will build up to. Shelly made a note that although we find objections to this or any view, this doesn’t mean that they aren’t worth considering, and since we can find objections about any view there isn’t a “perfect” answer to these questions. Our goal is to understand the flaws and strengths of the arguments and gain a better understanding of morality and ethics through this.

Mill does not own Utilitarianism, he merely wrote arguments to support it. Mill’s words on the matter are not God’s words on it (the same holds for Kant and *Kantianism*), so it’s important not to take this perspective.

Philosophy isn’t like physics; this isn’t a “learn it or fail” kind of course. You will need to decide for yourself what the right answers are and justify why you think that. There aren’t any *knock-down* arguments in Philosophy but this doesn’t mean that there aren’t right answers. To compare to Theology, there are disagreements about whether or not god exists but that doesn’t mean there is a truth value to that statement. Shelly wants us to know that he has bias himself and so do we, so it’s important to be aware of this. He’ll try to keep his cards close to his chest, so to speak.

**Problem 3.1.** If we believe something but there isn’t a set of arguments which will persuade any rational person of it, does that mean that our belief comes from a source outside of reason?

Shelly used the analogy of a jury deliberating to illustrate how there can be rational conflicts.

#### 3.1 Why have morality?

There are many possible answers. One possible answer is that different actions have different results, and we are not indifferent to the results of actions. We care about what happens in the world. This makes possible the thought that one should consider the re-

sults of one's actions; in a moral sense, we should consider whether the results are good or bad when deciding what one ought to do. Shelly calls this idea *Resultism*.

**Definition** (Resultism). The results of an action are morally relevant to determine the moral value of the action.

*Resultism*

The modest statement of this is that the results are *one of the things* which are relevant to this question; the strong statement says that results are the *only* things one should consider. This **bold resultism** is an agreement in Utilitarianism. Pretty much every moral theory has the modest version at least.

**Example 3.1** (Weak Resultism). If you read the paper and saw that many people died, almost everyone would agree that the results of that earthquake were bad (having been given no other information).

In order to use the results as a basis for morality, we need some way to measure the goodness or badness of results. This forms what's called a **theory of the good**. This lets us rank outcomes from better to worse. Here are some elements we may want to include in such a theory (intuitively plausible ideas, acceptable *prima facie*):

- What are the short term results?
- What are the long term results?
- What are the effects on *everybody*? (constats with the idea that morality should only consider how one's actions affect oneself, known as *ethical egoism*)

**Example 3.2** (Qualities we should care about?).

1. Is there a difference between planting a bomb in a school yard which will go off today, tomorrow, next year, in 100 years? Would any length of time be acceptable? [Long Term Results]
2. If we store nuclear waste in a cavern in Utah which won't leak for a thousand years, is that still acceptable if it kills 10,000 Utahans *eventually*? [Long Term Results]
3. Can we put our nuclear waste in Kenya so that no Americans are killed? [Who is affected]

The various flavors of ethical egoism:

1. Rational egoism: Asking “what should I do” as a matter of rationality; this is the in-house religion of the department of economics.
2. Ethical egoism: Should I consider how my actions affect someone else, or just consider what is good for me.
3. Psychological egoism: Our actions are entirely determined by what we think is good for ourselves individually.

You don’t need to believe one to believe the other. Bentham believed rational but not ethical.

**Example 3.3.** Suppose there is a lever which you can pull to get a candy bar. However, when you pull the lever 10,000 people in France are electrocuted. For an ethical egoist, this isn’t a moral dilemma: you should just pull the lever.

## 4 Thursday, January 24

“Ice cream is valuable because of the pleasurable sensations.”

---

Shelly

“I just don’t know what to say about masochists.”

---

Shelly

### 4.1 *What is the correct theory of the good?*

Generally we want to have a theory which is impartial and considers everyone equally. Sometimes we don’t want to single out some people over others. But that still doesn’t tell us how to pick one set of results over another.

**Proposition 4.1.** *Let’s measure how well-off people are after the different results. The modest way says that this is a relevant factor, the bolder way says it is the only thing that’s relevant.*

**Definition** (Welfarism). The view that people’s welfare is the only relevant criterion in determining which actions are morally better than another.

*Welfarism*

This still doesn’t give us a good account of well-being, and it doesn’t tell us how to aggregate individual well-being to understand how good or bad an outcome is as a whole. We have to answer both of these questions. Should it be based on material or personal goods, like money or vacations or respect or sex or one’s books to be published? In turn we need to understand why each of these good (or bad) things is valuable in its own right. It seems like the root of all these things is “pleasure” — everything else is an instrument to acquire pleasure.

**Problem 4.1.** Can pleasure be both inherently good and instrumentally good?  
Yes, cf. reproduction of genes through pleasurable activities.

**Example 4.1.** The pain you get at the dentist is intrinsically bad (because pain is bad) but instrumentally good (since it helps you avoid further pain and live a healthier life).



**Problem 4.2.** Can something be both intrinsically good and intrinsically bad? What about masochism?

**Problem 4.3.** How do we know if our list of inherently good and bad things complete? Are all pleasures equally good? Are all pains equally bad?

For this first question, Welfarism contends that pleasure and pain are the only good and bad things on that list.

**Definition** (Hedonism). The view that the value of a life is the sum difference of pleasures and pains.

*Hedonism*

**Example 4.2** (Painless Death). If your life is more good than bad, death would probably be bad. If your life is worse than it is good, death might be good. If you believe that the afterlife is amazing then death wouldn't really be a problem.

**Problem 4.4.** Do Hedonists consider the pleasure of other people?

**Definition** (Quantitative Hedonists). A hedonist who wants to measure the amount of pain or pleasure in an experience to rank them.

*Quantitative Hedonists*

## 5 Tuesday, January 29

“I’ve never been a pig.”

---

Shelly

Mill differentiates himself from Bentham by accepting that the quality of pleasures is important, even more so than the quantity of the pleasures. This leads to a hierarchy of pleasures which we should prefer. The social and intellectual and spiritual and aesthetic pleasures are ‘higher’ than the more base pleasures. In order to distinguish which are better than others, Mill proposes that if we ask those who understand and experience both and a decided majority support one over the other then that is the higher good.

**Problem 5.1.** How do we actually know that this is the answer people would give? Mill seems to privilege these particular ‘highbrow pleasures’ but then *justify* that choice by saying that ‘people will naturally choose that.’ But what if people don’t choose that — would Mill still hold those higher virtues above the lower ones?

**Problem 5.2.** If we accepted Qualitative Hedonism, wouldn’t that imply that we should force people to do these higher pleasures over the lower ones?

*Response.* Right now we’re discussing a theory of well-being, not a theory of what one ought to do. We might accept that someone might be happier if we forced them to do certain things, but we don’t say that we must force people to be happier, only recognize that this is something which would make them happier. Furthermore, *how* we experience the world is just as important as the thing we perceive; if we don’t like classical music then we aren’t experiencing any pleasure regardless of the supposed benefit of seeing it. ■

**Proposition 5.1.** *In order to understand and appreciate the higher virtues, we must train them to appreciate them through education.*

**Problem 5.3.** Five-year-olds seem to be a lot happier than adults.

*Response.* Would you really choose to be a five-year-old for the rest of your life? Most people (I think) wouldn’t. ■

**Problem 5.4.** Is getting the right mental state all that matters? Imagine that scientists can stimulate your brain directly to create experiences identical to those you would have while doing anything else. For example, you could fully experience what it would feel like to climb mount everest while inside the machine the whole time. But while you are in this machine, you don't know that you are in the machine since this knowledge would change your perspective. Do you want to be hooked up to the machine?

*Response (Personal).* I don't think so. ■

## 6 Thursday, January 31

“I went to the pound since Tuesday’s class and bought a stray cat. Brought it home, doused it in gasoline, and lit it on fire while it shrieked and died an agonizing death....and that seems wrong.”

---

Shelly

The results of the experience machine experiment tell us that most people don’t really buy Hedonism — there’s more to well-being than just the experience. There are many different theories which propose ideas about what the missing pieces are. One alternative is that we must have the things we want, and when we accomplish this our life is better than when we don’t. This theory explains why the experience machine fails; we want the real things, not merely the experience. These Preference Theories are the in-house philosophy of the department of economics.

It also seems like it ought to matter how welfare is aggregated. In some cases this is easy; if everybody is better off in world *A* than in world *B* it’s clear that *A* presents a better existence. But it’s often not so clear cut. Perhaps we ought to use the maximum happiness in the world, or perhaps we should use the average amount of happiness. There are scenarios which make either look bad. This has been going on for forty years and philosophers still don’t have a consensus. These theories all require us to have interpersonal comparisons of happiness or welfare, which may not even be possible.

**Problem 6.1.** Do animals count in the welfare aggregation?

## 7 No Notes

## 8 Thursday, February 7

“They’re graduate students in philosophy, they won’t contribute *anything* to society.”

---

Shelly

### *Does Utilitarianism Give Plausible Answers?*

**Problem 8.1.** Why is murder wrong?

*Utilitarianism Response.* Because you had other options, and the option to murder someone decreased happiness/welfare. ■

**Problem 8.2.** Should we steal from the blind man in the subway?

*Utilitarianism Response.* No. ■

**Problem 8.3.** We are sailing a boat when we see someone drowning; we turn to rescue them, but then we notice five people drowning somewhere else. We don’t have time to save everyone. What do we do?

*Utilitarianism Response.* It’s better to save five than to save one. ■

What if we tweak the scenario? What if the one was about the find the cure for cancer, and the five were all graduate students in philosophy. What about the impacts of what these people contribute to society?

*Utilitarianism Response.* Those things matter. We can’t say that everything is equal, as before. ■

What if we aren’t sure of the outcomes? We can’t know everything, after all. What should a utilitarianism (or a consequentialist) do in light of this? What if you accidentally save Hitler from drowning?

**Example 8.1.** There's a distinction here in what we're asking: Did you do the right thing or are you a morally good person? Utilitarianism and consequentialism are concerned with actions, and so in and of themselves do not concern themselves with whether or not people are good. We can develop a utilitarian theory of this, but it is supplemental to the core of the philosophy.